

Project 1: Mercedes-Benz Greener Manufacturing

by SHAKTI NATH SAINI

```
In [1]: #####
''' Step 1: Import the required libraries'''
#####

import numpy as np
import pandas as pd

# for dimensionality reduction
from sklearn.decomposition import PCA
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: #####
''' Step 2: Read the data from train.csv'''
#####

df_train = pd.read_csv(r'C:\Users\Shakti\Documents\Python Scripts\Projects\Mercedes-Benz Greener Manufacturing\Dataset\train.csv')

# Let us understand the data

print('Size of training set: {} rows and {} columns'
      .format(*df_train.shape))

# print few rows and see how the data Looks Like

df_train.head()
```

Size of training set: 4209 rows and 378 columns

```
Out[3]:
```

| | ID | y | X0 | X1 | X2 | X3 | X4 | X5 | X6 | X8 | ... | X375 | X376 | X377 | X378 | X379 | X380 | X382 | X383 | X384 | X385 | |
|---|----|--------|----|----|----|----|----|----|----|----|-----|------|------|------|------|------|------|------|------|------|------|---|
| 0 | 0 | 130.81 | k | v | a | t | a | d | u | j | o | ... | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 6 | 88.53 | k | t | a | v | e | d | y | l | o | ... | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 7 | 76.26 | az | w | n | c | d | x | j | x | ... | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |
| 3 | 9 | 80.62 | az | t | n | f | d | x | l | e | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 4 | 13 | 78.02 | az | v | n | f | d | h | d | n | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |

5 rows x 378 columns

```
In [4]: #####
''' Step 3: Collect the Y values into an array'''
#####

# separate the y from the data as we will use this to learn as
# the prediction output

y_train = df_train['y'].values
```

```
In [5]: #####
''' Step 4: Understand the data types we have'''
#####

# iterate through all the columns which has X in the name of the column

cols = [c for c in df_train.columns if 'X' in c]
print('Number of features: {}'.format(len(cols)))

print('Feature types:')
df_train[cols].dtypes.value_counts()

Number of features: 376
Feature types:
```

```
Out[5]: int64    368
object      8
dtype: int64
```

```
In [6]: #####
''' Step 5: Count the data in each of the columns'''
#####

counts = [[], [], []]
for c in cols:
    typ = df_train[c].dtype
    uniq = len(np.unique(df_train[c]))
    if uniq == 1:
        counts[0].append(c)
    elif uniq == 2 and typ == np.int64:
        counts[1].append(c)
    else:
        counts[2].append(c)

print('Constant features: {} Binary features: {} Categorical features: {}'.format(*[len(c) for c in counts]))
print('Constant features:', counts[0])
print('Categorical features:', counts[2])

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347']
Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']
```

```
In [7]: #####
''' Step 6: Read the test.csv data'''
#####

df_test = pd.read_csv(r'C:\Users\Shakti\Documents\Python Scripts\Projects\Mercedes-Benz Greener Manufacturing\Dataset\test.csv')

# remove columns ID and Y from the data as they are not used for learning

usable_columns = list(set(df_train.columns) - set(['ID', 'y']))
y_train = df_train['y'].values
id_test = df_test['ID'].values

x_train = df_train[usable_columns]
x_test = df_test[usable_columns]
```

```
In [8]: #####
''' Step 7: If for any column(s), the variance is equal to zero,
then you need to remove those variable(s).
Apply label encoder'''
#####

for column in usable_columns:
    cardinality = len(np.unique(x_train[column]))
    if cardinality == 1:
        x_train.drop(column, axis=1) # Column with only one
        # value is useless so we drop it
        x_test.drop(column, axis=1)

    if cardinality > 2: # Column is categorical
        mapper = lambda x: sum([ord(digit) for digit in x])
        x_train[column] = x_train[column].apply(mapper)
        x_test[column] = x_test[column].apply(mapper)

x_train.head()
```

Out[8]:

| | X370 | X177 | X151 | X346 | X28 | X101 | X368 | X205 | X227 | X361 | ... | X30 | X165 | X283 | X138 | X231 | X191 | X213 | X114 | X258 | X226 |
|---|------|------|------|------|-----|------|------|------|------|------|-----|-----|------|------|------|------|------|------|------|------|------|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1 | ... | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 0 | 1 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |

5 rows × 376 columns

```
In [9]: #####
''' Step 7: Check for null and unique values for test and train sets'''
#####

def check_missing_values(df):
    if df.isnull().any().any():
        print("There are missing values in the dataframe")
    else:
        print("There are no missing values in the dataframe")

check_missing_values(x_train)
check_missing_values(x_test)

There are no missing values in the dataframe
There are no missing values in the dataframe
```

```
In [10]: #####
''' Step 9: Make sure the data is now changed into numericals'''
#####

print('Feature types:')
x_train[cols].dtypes.value_counts()
```

Out[10]: int64 376
dtype: int64

```
In [11]: #####
''' Step10: Perform dimensionality reduction
Linear dimensionality reduction using Singular Value Decomposition of
the data to project it to a lower dimensional space.'''
#####

n_comp = 12
pca = PCA(n_components=n_comp, random_state=420)
pca2_results_train = pca.fit_transform(x_train)
pca2_results_test = pca.transform(x_test)
```

```
In [12]: #####
''' Step 11: Training using xgboost'''
#####

import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split

x_train, x_valid, y_train, y_valid = train_test_split(
    pca2_results_train,
    y_train, test_size=0.2,
    random_state=4242)

d_train = xgb.DMatrix(x_train, label=y_train)
d_valid = xgb.DMatrix(x_valid, label=y_valid)
#d_test = xgb.DMatrix(x_test)
d_test = xgb.DMatrix(pca2_results_test)

params = {}
params['objective'] = 'reg:linear'
params['eta'] = 0.02
params['max_depth'] = 4

def xgb_r2_score(preds, dtrain):
    labels = dtrain.get_label()
    return 'r2', r2_score(labels, preds)

watchlist = [(d_train, 'train'), (d_valid, 'valid')]

clf = xgb.train(params, d_train,
               1000, watchlist, early_stopping_rounds=50,
               feval=xgb_r2_score, maximize=True, verbose_eval=10)
```

[21:43:48] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/objective/regression_obj.cu:171: reg:linear is now deprecated in favor of reg:squarederror.

| | | | | |
|-------|---------------------|---------------------|---------------------|---------------------|
| [0] | train-rmse:99.14835 | train-r2: -58.35295 | valid-rmse:98.26297 | valid-r2: -67.63754 |
| [10] | train-rmse:81.27651 | train-r2: -38.88428 | valid-rmse:80.36433 | valid-r2: -44.91014 |
| [20] | train-rmse:66.71610 | train-r2: -25.87403 | valid-rmse:65.77334 | valid-r2: -29.75260 |
| [30] | train-rmse:54.86911 | train-r2: -17.17722 | valid-rmse:53.89136 | valid-r2: -19.64525 |
| [40] | train-rmse:45.24709 | train-r2: -11.36097 | valid-rmse:44.22322 | valid-r2: -12.90218 |
| [50] | train-rmse:37.44853 | train-r2: -7.46723 | valid-rmse:36.37614 | valid-r2: -8.40622 |
| [60] | train-rmse:31.14585 | train-r2: -4.85695 | valid-rmse:30.02252 | valid-r2: -5.40732 |
| [70] | train-rmse:26.08417 | train-r2: -3.10795 | valid-rmse:24.91497 | valid-r2: -3.41268 |
| [80] | train-rmse:22.04313 | train-r2: -1.93371 | valid-rmse:20.83055 | valid-r2: -2.08449 |
| [90] | train-rmse:18.84671 | train-r2: -1.14458 | valid-rmse:17.59679 | valid-r2: -1.20115 |
| [100] | train-rmse:16.33278 | train-r2: -0.61062 | valid-rmse:15.07888 | valid-r2: -0.61629 |
| [110] | train-rmse:14.39756 | train-r2: -0.25155 | valid-rmse:13.14822 | valid-r2: -0.22890 |
| [120] | train-rmse:12.92783 | train-r2: -0.00907 | valid-rmse:11.69290 | valid-r2: 0.02809 |
| [130] | train-rmse:11.80680 | train-r2: 0.15834 | valid-rmse:10.61692 | valid-r2: 0.19873 |
| [140] | train-rmse:10.98472 | train-r2: 0.27147 | valid-rmse:9.85685 | valid-r2: 0.30935 |
| [150] | train-rmse:10.37312 | train-r2: 0.35033 | valid-rmse:9.32088 | valid-r2: 0.38242 |
| [160] | train-rmse:9.91865 | train-r2: 0.40601 | valid-rmse:8.95928 | valid-r2: 0.42940 |
| [170] | train-rmse:9.58084 | train-r2: 0.44578 | valid-rmse:8.71328 | valid-r2: 0.46031 |
| [180] | train-rmse:9.33902 | train-r2: 0.47341 | valid-rmse:8.55257 | valid-r2: 0.48003 |
| [190] | train-rmse:9.15109 | train-r2: 0.49439 | valid-rmse:8.44777 | valid-r2: 0.49270 |
| [200] | train-rmse:8.98817 | train-r2: 0.51006 | valid-rmse:8.38791 | valid-r2: 0.49986 |
| [210] | train-rmse:8.90467 | train-r2: 0.52125 | valid-rmse:8.34975 | valid-r2: 0.50440 |
| [220] | train-rmse:8.82894 | train-r2: 0.52936 | valid-rmse:8.31976 | valid-r2: 0.50796 |
| [230] | train-rmse:8.76702 | train-r2: 0.53594 | valid-rmse:8.30397 | valid-r2: 0.50982 |
| [240] | train-rmse:8.72161 | train-r2: 0.54073 | valid-rmse:8.29916 | valid-r2: 0.51039 |
| [250] | train-rmse:8.67967 | train-r2: 0.54514 | valid-rmse:8.29209 | valid-r2: 0.51122 |
| [260] | train-rmse:8.64575 | train-r2: 0.54869 | valid-rmse:8.28976 | valid-r2: 0.51150 |
| [270] | train-rmse:8.61264 | train-r2: 0.55214 | valid-rmse:8.28928 | valid-r2: 0.51155 |
| [280] | train-rmse:8.58157 | train-r2: 0.55536 | valid-rmse:8.28961 | valid-r2: 0.51152 |
| [290] | train-rmse:8.55609 | train-r2: 0.55800 | valid-rmse:8.28593 | valid-r2: 0.51195 |
| [300] | train-rmse:8.53319 | train-r2: 0.56036 | valid-rmse:8.28463 | valid-r2: 0.51210 |
| [310] | train-rmse:8.50275 | train-r2: 0.56349 | valid-rmse:8.28408 | valid-r2: 0.51217 |
| [320] | train-rmse:8.48186 | train-r2: 0.56564 | valid-rmse:8.28211 | valid-r2: 0.51240 |
| [330] | train-rmse:8.45378 | train-r2: 0.56851 | valid-rmse:8.27860 | valid-r2: 0.51281 |
| [340] | train-rmse:8.42514 | train-r2: 0.57143 | valid-rmse:8.27825 | valid-r2: 0.51285 |
| [350] | train-rmse:8.40047 | train-r2: 0.57393 | valid-rmse:8.28035 | valid-r2: 0.51261 |
| [360] | train-rmse:8.37167 | train-r2: 0.57685 | valid-rmse:8.27932 | valid-r2: 0.51273 |
| [370] | train-rmse:8.35221 | train-r2: 0.57881 | valid-rmse:8.27776 | valid-r2: 0.51291 |
| [380] | train-rmse:8.33159 | train-r2: 0.58089 | valid-rmse:8.27285 | valid-r2: 0.51349 |
| [390] | train-rmse:8.30856 | train-r2: 0.58320 | valid-rmse:8.27193 | valid-r2: 0.51360 |
| [400] | train-rmse:8.28113 | train-r2: 0.58595 | valid-rmse:8.26742 | valid-r2: 0.51413 |
| [410] | train-rmse:8.25393 | train-r2: 0.58867 | valid-rmse:8.26621 | valid-r2: 0.51427 |
| [420] | train-rmse:8.22845 | train-r2: 0.59120 | valid-rmse:8.26609 | valid-r2: 0.51428 |
| [430] | train-rmse:8.20038 | train-r2: 0.59399 | valid-rmse:8.26503 | valid-r2: 0.51441 |
| [440] | train-rmse:8.17690 | train-r2: 0.59631 | valid-rmse:8.26380 | valid-r2: 0.51455 |
| [450] | train-rmse:8.15205 | train-r2: 0.59876 | valid-rmse:8.26049 | valid-r2: 0.51494 |
| [460] | train-rmse:8.12528 | train-r2: 0.60139 | valid-rmse:8.26295 | valid-r2: 0.51465 |
| [470] | train-rmse:8.10488 | train-r2: 0.60339 | valid-rmse:8.26308 | valid-r2: 0.51464 |
| [480] | train-rmse:8.08610 | train-r2: 0.60523 | valid-rmse:8.26411 | valid-r2: 0.51452 |
| [490] | train-rmse:8.06069 | train-r2: 0.60770 | valid-rmse:8.26696 | valid-r2: 0.51418 |
| [499] | train-rmse:8.04085 | train-r2: 0.60963 | valid-rmse:8.26537 | valid-r2: 0.51437 |

```
In [13]: #####
''' Step 12: Predict your test_df values using xgboost
'''
#####

p_test = clf.predict(d_test)

sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('xgb.csv', index=False)

sub.head()
```

Out[13]:

| | ID | y |
|---|----|------------|
| 0 | 1 | 82.949509 |
| 1 | 2 | 97.652306 |
| 2 | 3 | 83.115051 |
| 3 | 4 | 76.980057 |
| 4 | 5 | 112.474335 |

```
In [14]: #####
'''
'''
#####
```

Out[14]:

| | | |
|--|-----|--|
| | End | |
|--|-----|--|

